

Digital Signals and Image Processing

Semester 2, 2024

Dr Maryam Haghighat

Tarang Janawalkar

This work is licensed under a Creative Commons
“Attribution-NonCommercial-ShareAlike 4.0 International” license.



Contents

Contents	1
1 Digital Image Processing	3
1.1 Elements of Visual Perception	3
1.1.1 Image Formation	3
1.1.2 Brightness Adaptation and Discrimination	3
1.2 Image Generation Components	3
1.2.1 Electromagnetic Spectrum	3
1.2.2 Image Sensors	4
1.2.3 Human Perception	4
1.3 Image Sensing and Acquisition	4
1.3.1 Image Sensing Modalities	4
1.3.2 Image Formation	5
1.4 Image Sampling and Quantisation	5
1.4.1 Dynamic Range and Contrast	6
1.4.2 Spatial and Intensity Resolution	6
1.4.3 Image Interpolation	6
1.5 Relationships Between Pixels	7
1.5.1 Neighbours of a Pixel	7
1.5.2 Adjacency and Connectivity	7
1.5.3 Connectivity	8
1.6 Distance Metrics	8
1.7 Mathematical Operations	9
1.7.1 Element-wise Operations	9
1.7.2 Linear Operations	9
1.7.3 Spatial Operations	9
2 Intensity Transformations & Spatial Filters	9
2.1 Intensity Transformations	9
2.1.1 Identity Transformation	9
2.1.2 Negative Transformation	10
2.1.3 Logarithmic Transformation	10
2.1.4 Power-Law (Gamma) Transformation	10
2.1.5 Piecewise-Linear Transformation	10
2.2 Histogram Processing	11
2.2.1 Histogram Equalisation	11
2.2.2 Histogram Matching	11
2.3 Spatial Filtering	11
2.3.1 Correlation	12
2.3.2 Convolution	12
2.3.3 Padding	12
2.3.4 Averaging Filters	12
2.3.5 Smoothing Linear Filters	13
2.3.6 Order-Statistic Non-Linear Filters	13

2.3.7	Sharpening Filters	13
2.3.8	Unsharp Masking and High-Boost Filtering	14
3	Filtering in the Frequency Domain	14
3.1	2D Fourier Transform	14
3.2	Filtering in the Frequency Domain	15
3.3	Image Smoothing	15
3.3.1	Ideal Low-Pass Filters	15
3.3.2	Butterworth Low-Pass Filters	16
3.3.3	Gaussian Low-Pass Filters	16
3.4	Image Sharpening	16
3.4.1	Ideal High-Pass Filters	16
3.4.2	Butterworth High-Pass Filters	16
3.4.3	Gaussian High-Pass Filters	16
3.4.4	Laplacian Filters	17
3.5	Selective Filtering	17
3.5.1	Butterworth Bandreject Filters	17
3.5.2	Gaussian Bandreject Filters	18
3.5.3	Notch Filters	18
4	Colour	18
4.1	Colour Fundamentals	18
4.1.1	Characterising Colour	19
4.2	Colour Models	19
4.2.1	RGB Colour Model	19
4.2.2	CMY(K) Colour Model	20
4.2.3	HSI Colour Model	20
4.3	Colour Image Acquisition	21
4.3.1	Pixel Depth	21
4.4	Colour Image Processing	21
4.4.1	Colour Intensity Transformations	21
4.4.2	Colour Complements	22
4.4.3	Colour Slicing	22
4.4.4	Colour Tonal Transformations	22
4.4.5	Colour Balancing (White Balancing)	22
4.4.6	Colour Histogram Equalisation	22
4.4.7	Colour Image Smoothing and Sharpening	22
4.4.8	Colour Image Segmentation	23

1 Digital Image Processing

Digital image processing is the processing of images on a digital system using algorithms. A digital image is a binary representation of visual data that is composed of a finite number of elements, each with a particular location and value.

Image processing methods can be divided into two categories:

- Methods where the input and output are images.
- Methods where the input is an image and the output is some information extracted from the image.

1.1 Elements of Visual Perception

The human visual system has influenced and contributed to many advancements in image processing. The human eye has light receptors called rods and cones. Humans have around 6 to 7 million cones in each eye that are highly sensitive to colour and fine details. On the other hand, there are a total of 75 to 150 million rods across both eyes, that are sensitive to low levels of illumination.

1.1.1 Image Formation

Photo camera lenses are fixed in focal length and they focus at various distances by varying the distance between the lens and imaging plane (film/chip). The human eye works in the opposite way, where the distance between the lens and the imaging plane (retina) is fixed, but the focal length for focus is varied by changing the shape of the lens.

1.1.2 Brightness Adaptation and Discrimination

The human eye is capable of discriminating between a wide range of intensity levels. This range of light intensity levels is on the order of 10^{10} .

1.2 Image Generation Components

There are three components to image generation:

- **Object:** The object being imaged.
- **Energy Source:** The source of energy that illuminates the object.
- **Sensor:** The sensor that detects the energy reflected from the object.

Depending on the properties of the energy source and object material and geometry, the emitted energy can be reflected, transmitted, or absorbed.

1.2.1 Electromagnetic Spectrum

The main source of energy for imaging is electromagnetic (EM) radiation. EM radiation consists of propagating sinusoidal waves characterised by their oscillating frequency f . Using Planck's equation, the energy of a photon can be calculated as

$$E = hf$$

where $h = 6.626\,070\,15 \times 10^{-34}$ J s is Planck's constant and f is the frequency of the wave. Given the speed of light $c = 299\,792\,458$ m s⁻¹, we can also calculate the energy using

$$E = \frac{hc}{\lambda}.$$

The EM spectrum is divided into regions based on the frequency of the waves. The visible light spectrum is a small part of the EM spectrum that is visible to the human eye. The visible light spectrum ranges from 380 nm to 700 nm. Earth's atmosphere also blocks certain parts of the EM spectrum, such as short wavelength UV and X-rays.

Perceived colour (hue) is related to the wavelength of light, while the **brightness** is related to the intensity of the radiation.

1.2.2 Image Sensors

Image sensors capture a specific range of the EM spectrum, for example:

- RGB sensors capture the visible light spectrum.
- Infrared sensors capture the infrared spectrum.
- X-ray sensors capture the X-ray spectrum.
- Ultraviolet sensors capture the ultraviolet spectrum.

1.2.3 Human Perception

Human perception is context-dependent. Perceived intensity around regions of discontinuous intensity appear to undershoot and overshoot around the boundary (see Mach band effect). The eye can also fill in non-existing information and wrongly perceive geometrical properties of objects. To produce a powerful vision system, we need both a powerful image sensor and image processor to extract useful information from an image.

1.3 Image Sensing and Acquisition

Image sensing is the process of transforming illuminated energy into a digital image. The process involves the following steps:

1. Convert the illuminated energy into an electrical signal.
2. Digitize the electrical signal to obtain a digital image.

1.3.1 Image Sensing Modalities

Image sensing is done using three principal modalities:

- **Single Sensing Element:** A single sensor that captures the energy. For example, a photodiode. To generate 2D images, the sensor must be appropriately displaced in the x and y directions.

- **Line Sensor:** A sensor that captures energy along a line. To generate 2D images, the sensor must be displaced in the direction perpendicular to the line.
- **Array Sensor:** A sensor that captures energy in a 2D array. The sensor is divided into rows and columns, with each element capturing energy at a specific location. A typical arrangement is the **CCD** (Charge Coupled Device) sensor.

1.3.2 Image Formation

Let us denote the intensity of a monochrome image by the 2-dimensional function

$$\ell = f(x, y)$$

where x and y represent the spatial coordinates captured by the sensor, and f is a scalar function of the intensity of the energy radiated by a physical source. As such, this function is non-negative and finite:

$$0 \leq \ell \leq \infty.$$

f is characterised by two components:

- **Illumination:** The amount of source illumination incident on the scene $i(x, y)$. Here $0 \leq i(x, y) < \infty$.
- **Reflectance:** The amount of illumination reflected by the objects on the scene $r(x, y)$. Here $0 \leq r(x, y) \leq 1$.

Therefore, we can describe the image formation process as

$$f(x, y) = i(x, y) r(x, y),$$

where $r = 0$ implies total absorption, while $r = 1$ implies total reflectance. For monochrome images, we can define the minimum and maximum intensity values as L_{\min} and L_{\max} , respectively, where

$$L_{\min} \leq \ell \leq L_{\max}, \quad L_{\min} = i_{\min} r_{\min}, \quad L_{\max} = i_{\max} r_{\max}.$$

The range of intensity values $[L_{\min}, L_{\max}]$ is called the **intensity/gray scale**. Commonly, this interval is transformed to the interval $[0, L - 1]$, where L is the number of intensity levels.

1.4 Image Sampling and Quantisation

Image sampling is the process of sampling discrete points in a continuous image. Regardless of the sensor arrangement, the image is sampled at a fixed rate in the x and y directions and the resulting points are called **pixels**. These pixels are stored in an array that is $M \times N$ in size, where M is the number of rows and N is the number of columns, arranged as shown below:

$$\mathbf{X} = \begin{bmatrix} x_{00} & x_{01} & \cdots & x_{0,N-1} \\ x_{10} & x_{11} & \cdots & x_{1,N-1} \\ \vdots & \vdots & \ddots & \vdots \\ x_{M-1,0} & x_{M-1,1} & \cdots & x_{M-1,N-1} \end{bmatrix}$$

Image quantisation is the process of converting the continuous intensity values of an image to discrete values. The number of intensity levels L is determined by the number of bits used to represent each pixel. The number of intensity levels is given by

$$L = 2^k,$$

where k is the bit depth. Thus the quantised intensity values will range from $[0, L - 1]$. The **quality** of an image is determined by the number of discrete intensity levels used in both sampling and quantisation.

1.4.1 Dynamic Range and Contrast

The **dynamic range** of an image is the ratio of the maximum intensity value to the minimum intensity value:

$$\text{Dynamic Range} = \frac{L_{\max}}{L_{\min}} = \frac{i_{\max}r_{\max}}{i_{\min}r_{\min}}.$$

The upper limit is determined by the sensor's saturation level, while the lower limit is determined by the sensor's noise level. The **contrast** of an image is the difference in intensity between the brightest and darkest regions of the image:

$$\text{Contrast} = L_{\max} - L_{\min} = i_{\max}r_{\max} - i_{\min}r_{\min}.$$

- A **high** dynamic range implies a large difference between the brightest and darkest regions of the image, and therefore high contrast.
- A **low** dynamic range implies a small difference between the brightest and darkest regions of the image, and therefore low contrast.

1.4.2 Spatial and Intensity Resolution

The **spatial resolution** of an image is a measure of the smallest discernible detail in an image, measured as the number of pixels per unit area (dots per inch). In some images, sampling an image at a low rate can result in **aliasing**, where high-frequency components are incorrectly represented as low-frequency components (see the Moiré pattern).

The **intensity resolution** of an image is the smallest discernible change in intensity level, which is related to the number of intensity levels used to represent the image, for example, 8-bit and 10-bit images. Choosing a low number of intensity levels can result in quantisation noise, where intensity levels are incorrectly represented.

1.4.3 Image Interpolation

Image interpolation is the process of estimating the intensity values of pixels between the sampled points. Interpolation is used to increase and decrease the resolution of an image for resampling and resizing. Common interpolation methods include:

- **Nearest Neighbour:** The intensity value of the nearest pixel is used to estimate the intensity value of the pixel.

- **Bilinear:** The intensity value of the nearest four pixels is used to estimate the intensity value of the pixel.
- **Bicubic:** The intensity value of the nearest sixteen pixels is used to estimate the intensity value of the pixel.

1.5 Relationships Between Pixels

The following sections will define some common sets that are used to describe relationships between pixels in an image.

1.5.1 Neighbours of a Pixel

The **neighbours** of a pixel $\mathbf{p} = (x, y)$ are the pixels that are adjacent to \mathbf{p} .

- The **4-neighbours** of \mathbf{p} are defined as the pixels that are adjacent to \mathbf{p} in the **cardinal** directions:

$$\mathbf{N}_4(\mathbf{p}) = \{(x, y-1), (x-1, y), (x+1, y), (x, y+1)\},$$

- The **diagonal-neighbours** of \mathbf{p} are defined as the pixels that are adjacent to \mathbf{p} in the **diagonal** directions:

$$\mathbf{N}_D(\mathbf{p}) = \{(x-1, y-1), (x+1, y-1), (x-1, y+1), (x+1, y+1)\},$$

- The **8-neighbours** of \mathbf{p} are defined as the pixels that are adjacent to \mathbf{p} in both the **cardinal** and **diagonal** directions:

$$\mathbf{N}_8(\mathbf{p}) = \mathbf{N}_4(\mathbf{p}) \cup \mathbf{N}_D(\mathbf{p}).$$

1.5.2 Adjacency and Connectivity

Two pixels \mathbf{p} and \mathbf{q} are **adjacent** if they are neighbours and their intensity values are similar or belong to the same set of values V based on a threshold. This can occur in one of three ways:

- **4-Adjacency** when $\mathbf{q} \in \mathbf{N}_4(\mathbf{p})$.
- **8-Adjacency** when $\mathbf{q} \in \mathbf{N}_8(\mathbf{p})$.
- **M-Adjacency** when $\mathbf{q} \in \mathbf{N}_4(\mathbf{p})$ or $\mathbf{q} \in \mathbf{N}_D(\mathbf{p})$ and $\mathbf{N}_4(\mathbf{p}) \cap \mathbf{N}_4(\mathbf{q}) = \emptyset$. This statement avoids double-counting an adjacency when \mathbf{q} is a diagonal neighbour of \mathbf{p} while another cardinal neighbour exists between \mathbf{p} and \mathbf{q} . Consider the binary example with $\mathbf{p}, \mathbf{q}, \mathbf{r} \in V = \{1\}$:

$$\begin{bmatrix} 0 & \mathbf{r} & \mathbf{q} \\ 0 & \mathbf{p} & 0 \\ 0 & 0 & 0 \end{bmatrix}.$$

Here we do not consider \mathbf{p} and \mathbf{q} to be M-adjacent, as \mathbf{p} and \mathbf{q} are already adjacent through \mathbf{r} .

A **path** (or curve) between two pixels \mathbf{p} and \mathbf{q} is a sequence of $n+1$ pixels $(\mathbf{p}_0, \mathbf{p}_1, \dots, \mathbf{p}_n)$ such that $\mathbf{p}_0 = \mathbf{p}$ and $\mathbf{p}_n = \mathbf{q}$, where \mathbf{p}_i is adjacent to \mathbf{p}_{i+1} for $i = 0, 1, 2, \dots, n-1$. This path is said to have length n , and is *closed* if $\mathbf{p}_0 = \mathbf{p}_n$.

1.5.3 Connectivity

Consider a subset of pixels in an image S with $\mathbf{p}, \mathbf{q} \in S$.

- \mathbf{p} and \mathbf{q} are **connected** if there exists a path between \mathbf{p} and \mathbf{q} such that all pixels in the path are in S .
- The set of pixels connected to \mathbf{p} in S form a **connected component**.
- If S only consists of *one* connected component, then S is said to be a **connected set** and is called a **region** R .

Two regions R_1 and R_2 are adjacent if their union forms a connected set, i.e., another region. Regions that are not adjacent are said to be **disjoint**.

1.6 Distance Metrics

The **distance** between two pixels \mathbf{p} and \mathbf{q} can be measured using a variety of metrics. The most common metrics are:

- **Euclidean distance:**

$$d(\mathbf{p}, \mathbf{q}) = \sqrt{(x_p - x_q)^2 + (y_p - y_q)^2}.$$

- **Manhattan (D4) distance** (or city-block distance):

$$d_4(\mathbf{p}, \mathbf{q}) = |x_p - x_q| + |y_p - y_q|.$$

- **Chessboard (D8) distance** (or maximum distance):

$$d_8(\mathbf{p}, \mathbf{q}) = \max(|x_p - x_q|, |y_p - y_q|).$$

In general, any metric d that satisfies the following properties is a **distance metric**:

- The distance from a point to itself is zero:

$$d(\mathbf{p}, \mathbf{p}) = 0.$$

- **Positivity:** The distance between two distinct points is always positive:

$$d(\mathbf{p}, \mathbf{q}) > 0 : \mathbf{p} \neq \mathbf{q}.$$

- **Symmetry:** The distance between two points is always the same regardless of ordering:

$$d(\mathbf{p}, \mathbf{q}) = d(\mathbf{q}, \mathbf{p}).$$

- The Triangle Inequality is satisfied:

$$d(\mathbf{p}, \mathbf{q}) \leq d(\mathbf{p}, \mathbf{r}) + d(\mathbf{r}, \mathbf{q}).$$

for all $\mathbf{p}, \mathbf{q}, \mathbf{r}$ in this metric space.

1.7 Mathematical Operations

1.7.1 Element-wise Operations

Given two images \mathbf{X} and \mathbf{Y} of equal dimensions, an element-wise operation is an operation that is applied to each pixel in the image. Suppose we wish to apply the binary operation \otimes on \mathbf{X} and \mathbf{Y} . The resulting image \mathbf{Z} is given by

$$\mathbf{Z} = \mathbf{X} \otimes \mathbf{Y} = \begin{bmatrix} x_{00} \otimes y_{00} & x_{01} \otimes y_{01} & \cdots & x_{0,N-1} \otimes y_{0,N-1} \\ x_{10} \otimes y_{10} & x_{11} \otimes y_{11} & \cdots & x_{1,N-1} \otimes y_{1,N-1} \\ \vdots & \vdots & \ddots & \vdots \\ x_{M-1,0} \otimes y_{M-1,0} & x_{M-1,1} \otimes y_{M-1,1} & \cdots & x_{M-1,N-1} \otimes y_{M-1,N-1} \end{bmatrix}$$

Here the operator \otimes can represent any binary operation, such as addition, subtraction, multiplication, or division.

1.7.2 Linear Operations

The operator H is a linear operation if it satisfies the following property:

$$H[a\mathbf{X} + b\mathbf{Y}] = aH[\mathbf{X}] + bH[\mathbf{Y}],$$

for images \mathbf{X} and \mathbf{Y} , where a and b are constants.

1.7.3 Spatial Operations

The operator T is a spatial operation if the output pixel value is determined by the values of pixels in the neighbourhood of any input pixel. T can be categorised as one of three types of spatial operations:

- **Pointwise:** The output pixel value is determined by the value of the input pixel only.
- **Neighbourhood:** The output pixel value is determined by the value of the input pixel and any neighbouring pixels.
- **Global:** The output pixel value is determined by the value of pixels in the entire image.
- **Geometric:** The output pixel value is determined by the spatial location of the pixel.

2 Intensity Transformations & Spatial Filters

2.1 Intensity Transformations

An intensity transform aims to modify the contrast of an image by changing the range of intensity values in that image. The following sections will define some common intensity transformations. In the following sections let r be the input intensity and s be the output intensity of an image.

2.1.1 Identity Transformation

The identity transformation is the simplest intensity transformation which does not alter an image. It is defined as

$$s = T(r) = r.$$

2.1.2 Negative Transformation

The image negation transformation is used to invert the intensity values of an image. It is defined as

$$s = T(r) = (L - 1) - r,$$

where L is the number of intensity levels in the image.

2.1.3 Logarithmic Transformation

The logarithmic transformation is used to enhance darker regions of an image by compressing brighter regions. It is defined as

$$s = T(r) = c \log(1 + r),$$

where c is a constant that scales the intensity values of the image.

2.1.4 Power-Law (Gamma) Transformation

The power-law transformation is used to correct the gamma of an image, either by enhancing or reducing dark or bright regions. It can be thought of as a generalisation of the logarithmic transformation. It is defined as

$$s = T(r) = cr^\gamma,$$

where c is a constant that scales the intensity values of the image and γ is the gamma value.

- When $\gamma < 1$, the transformation enhances the darker regions of the image, while compressing the brighter regions.
- When $\gamma > 1$, the transformation enhances the brighter regions of the image, while compressing the darker regions.

2.1.5 Piecewise-Linear Transformation

Piecewise-linear transformations are used to enhance the contrast of specific regions of an image. Some common piecewise-linear transformations include:

- **Contrast Stretching:** Enhances the contrast of an image by stretching the intensity values to the full range of intensity levels.
- **Intensity Level Slicing:** Enhances the contrast of specific regions of an image by setting the intensity values of other regions to zero or by leaving them unchanged.
- **Bit-Plane Slicing:** Highlights the contribution made to image appearance by specific bits in the image.

2.2 Histogram Processing

Histograms are used to visualise the distribution of intensity values in an image. Given an image $\mathbf{X} \in [0, L-1]^{M \times N}$, the histogram $h_{\mathbf{X}}(k)$ is defined as

$$h_{\mathbf{X}}(k) = n_k,$$

where n_k is the number of pixels in the image with intensity value k . If we normalise these values, we find the probability of obtaining a pixel with intensity value k :

$$p_{\mathbf{X}}(k) = \frac{n_k}{MN}.$$

It follows that

$$\sum_{k=0}^{L-1} p_{\mathbf{X}}(k) = 1.$$

2.2.1 Histogram Equalisation

Histogram equalisation is a method used to spread the most frequent intensity values in an image to the full range of intensity levels, thereby achieving a more uniform distribution of intensity values. To do this, we will use the following transformation that maps the cumulative distribution function of an input image \mathbf{X} to the cumulative distribution function of a uniform distribution:

$$s = T(r) = (L-1) \sum_{j=0}^r p_{\mathbf{X}}(j) = \frac{L-1}{MN} \sum_{j=0}^r n_j.$$

2.2.2 Histogram Matching

In some cases, we wish to match the histogram of \mathbf{X} to the histogram of another image \mathbf{Y} . To do so, consider the histogram equalisation of \mathbf{Y} with intensity values z :

$$s = G(z) = (L-1) \sum_{j=0}^z p_{\mathbf{Y}}(j) = \frac{L-1}{MN} \sum_{j=0}^z n_j.$$

Thus we have the mapping:

$$T : r \mapsto s \quad \text{and} \quad G : z \mapsto s.$$

As both T and G map to the same equalised space, we can define the transformation $z = H(r)$ that maps the histogram of \mathbf{X} to the histogram of \mathbf{Y} as

$$z = H(r) = G^{-1}(T(r)) = G^{-1}(s).$$

2.3 Spatial Filtering

Spatial filtering is the process of creating a new image by applying a mask (or kernel, template or window) to each pixel in an image. This new pixel value is determined by the intensity values of the pixels in the neighbourhood of the original pixel. The mask is defined as a $m \times n$ matrix \mathbf{W} with elements w_{ij} that represent the weights of the pixels in the neighbourhood of the pixel being processed. For convenience, m and n are typically odd integers.

The output pixel value is given by the weighted sum of the intensity values of the pixels in the neighbourhood of the pixel. This can be done using one of two operations:

- **Correlation:** The mask is shifted across the image and the weighted sum is calculated at each position.
- **Convolution:** The mask is first flipped horizontally and vertically before it is shifted across the image.

2.3.1 Correlation

The (i, j) th element of the correlation of an image \mathbf{X} with a mask \mathbf{W} is defined as

$$y_{ij} = \mathbf{W} \star \mathbf{X} = \sum_{s=-a}^a \sum_{t=-b}^b w_{st} x_{i+s, j+t}$$

for $a = \frac{m-1}{2}$ and $b = \frac{n-1}{2}$.

2.3.2 Convolution

The (i, j) th element of the convolution of an image \mathbf{X} with a mask \mathbf{W} is defined as

$$y_{ij} = \mathbf{W} * \mathbf{X} = \sum_{s=-a}^a \sum_{t=-b}^b w_{st} x_{i-s, j-t}$$

for $a = \frac{m-1}{2}$ and $b = \frac{n-1}{2}$.

2.3.3 Padding

For masks larger than 1×1 , the indices w_{st} will exceed the bounds of the image \mathbf{X} . To prevent this, the image is often padded with an additional border of pixels. Common padding methods include:

- **Zero Padding:** The border is padded with zeros.
- **Boundary Replication Padding:** The border is padded with the intensity values of the nearest pixel.
- **Reflection Padding:** The border is padded with the intensity values of reflected pixels (one pixel away from the border).

2.3.4 Averaging Filters

Averaging filters are used to reduce noise in an image by averaging the intensity values of the pixels in the neighbourhood of the pixel being processed. An averaging filter considers a continuous function of two variables, such as the multivariable Gaussian function:

$$w_{st} = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{s^2 + t^2}{2\sigma^2}\right).$$

2.3.5 Smoothing Linear Filters

Smoothing filters are used to reduce noise in an image by averaging the intensity values of the pixels in the neighbourhood of the pixel being processed. A general implementation for filtering an $M \times N$ image with a **weighted averaging filter** of size $m \times n$ is defined as:

$$y_{ij} = \frac{\sum_{s=-a}^a \sum_{t=-b}^b w_{st} x_{i+s, j+t}}{\sum_{s=-a}^a \sum_{t=-b}^b w_{st}},$$

2.3.6 Order-Statistic Non-Linear Filters

Order-statistic filters are used to reduce noise in an image by replacing the intensity value of a pixel with the median, maximum, or minimum intensity value of the pixels in the neighbourhood of the pixel being processed. Median filters have good noise-reduction capabilities with less smoothing and are used to remove impulse or salt-and-pepper noise.

2.3.7 Sharpening Filters

Sharpening filters are used to enhance edges and discontinuities in an image. One technique is to consider the Laplacian of the image $f(x, y)$:

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}.$$

For discrete images, we will use the first-order forward difference approximation:

$$\begin{aligned} \frac{\partial f}{\partial x} &= f(x+1, y) - f(x, y) \\ \frac{\partial f}{\partial y} &= f(x, y+1) - f(x, y), \end{aligned}$$

and the second-order central difference approximation:

$$\begin{aligned} \frac{\partial^2 f}{\partial x^2} &= f(x+1, y) - 2f(x, y) + f(x-1, y) \\ \frac{\partial^2 f}{\partial y^2} &= f(x, y+1) - 2f(x, y) + f(x, y-1). \end{aligned}$$

The Laplacian allows us to identify transitions in intensity values across an image by creating a filter with one of the following masks:

$$\mathbf{W} = \begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix} \quad \text{or} \quad \mathbf{W} = \begin{bmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{bmatrix},$$

where the first mask will sharpen the image, while the second mask will sharpen the negative of the image. In general, we can sharpen edges by adding the Laplacian to the original image:

$$g(x, y) = f(x, y) + c \nabla^2 f(x, y),$$

where $c = -1$ will sharpen the image, while $c = 1$ sharpen the negative of the image.

2.3.8 Unsharp Masking and High-Boost Filtering

Unsharp masking is a sharpening technique that enhances edges and discontinuities in an image by subtracting a blurred version of the image from the original image. The process takes the following steps:

1. Blur the original image $f(x, y)$.
2. Subtract the blurred image from the original image to obtain the mask

$$g_{\text{mask}}(x, y) = f(x, y) - \bar{f}(x, y),$$

where $\bar{f}(x, y)$ is the blurred image.

3. Add the mask to the original image to obtain a sharpened image:

$$g(x, y) = f(x, y) + k g_{\text{mask}}(x, y), \quad k > 0.$$

where

- $k = 1$ corresponds to unsharp masking.
- $k > 1$ corresponds to high-boost filtering.

3 Filtering in the Frequency Domain

3.1 2D Fourier Transform

The 2D Fourier transform of an image $f(x, y)$ ¹ is defined as

$$F(u, v) = \mathcal{F}\{f(x, y)\} = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x, y) e^{-j2\pi(ux+vy)} dx dy,$$

with the inverse 2D Fourier transform defined as:

$$f(x, y) = \mathcal{F}^{-1}\{F(u, v)\} = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} F(u, v) e^{j2\pi(ux+vy)} du dv.$$

When f is a discrete $M \times N$ image, the 2D discrete Fourier transform is defined as

$$F(u, v) = \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} f(x, y) e^{-j2\pi(ux/M+vy/N)},$$

with the inverse 2D discrete Fourier transform defined as

$$f(x, y) = \sum_{u=0}^{M-1} \sum_{v=0}^{N-1} F(u, v) e^{j2\pi(ux/M+vy/N)}.$$

Many properties of the Fourier transform in the 1D case extend to the 2D case, such as linearity, symmetry, and shift properties. Notably, the sampling theorem in 2D states that an image can be

¹ f must satisfy the Dirichlet conditions for the 2D Fourier transform to exist.

perfectly reconstructed from its samples if the sampling rate f_s is greater than twice the highest frequency component in the image:

$$f_s > 2f_{\max}.$$

In 2D, this is equivalent to:

$$\frac{1}{\Delta x} > 2u_{\max} \quad \text{and} \quad \frac{1}{\Delta y} > 2v_{\max},$$

where u_{\max} and v_{\max} are the maximum frequencies in the x and y directions, respectively. The frequency of spatial and frequency samples can also be related by the following equations:

$$\Delta u = \frac{1}{M \Delta x} \quad \text{and} \quad \Delta v = \frac{1}{N \Delta y}.$$

3.2 Filtering in the Frequency Domain

Filtering can also be performed in the frequency domain if the image contains noise that is more easily removed when visualising the magnitude and phase spectra of the image. In general, a filter in the frequency domain is defined as

$$g(x, y) = \mathcal{F}^{-1} \{H(u, v) F(u, v)\},$$

where $H(u, v)$ is a filter designed to act on the frequency spectrum of the image. Due to the lack of padding, the horizontal or vertical edges of the resulting image may contain black pixels, leading to inconsistent filtering. To avoid this, we can take the following steps:

1. For an image $f(x, y)$ of size $M \times N$, define padding parameters $P = 2M$ and $Q = 2N$.
2. Form a padded image $f_p(x, y)$ of size $P \times Q$ by appending the necessary number of zeros to f .
3. Multiply the padded image by $(-1)^{x+y}$ to centre the transform at the origin.
4. Compute the 2D Fourier transform of the transformed padded image.
5. Multiply the result by the symmetric filter $H(u, v)$ of size $P \times Q$ to form $G(u, v)$.
6. Obtain the processed image $g_p(x, y) = \Re \{ \mathcal{F}^{-1} \{G(u, v)\} \} (-1)^{x+y}$.
7. Obtain the final processed image $g(x, y)$ by cropping the processed image to the original size $M \times N$, which is in the top-left corner of the processed image.

3.3 Image Smoothing

3.3.1 Ideal Low-Pass Filters

An ideal low-pass filter passes all frequencies within a circle of radius D_0 from the origin and suppresses all frequencies outside of this circle, without attenuation:

$$H(u, v) = \begin{cases} 1 & \text{if } D(u, v) \leq D_0, \\ 0 & \text{if } D(u, v) > D_0, \end{cases}$$

Here $D(u, v)$ is the distance from the origin to the point (u, v) in the frequency domain:

$$D(u, v) = \sqrt{(u - P/2)^2 + (v - Q/2)^2}.$$

3.3.2 Butterworth Low-Pass Filters

A Butterworth low-pass filter is defined as

$$H(u, v) = \frac{1}{1 + (D(u, v)/D_0)^{2n}},$$

where n is the order of the filter. A higher-order filter has a sharper transition between the passband and the stopband, but also results in more rippling in the spatial domain.

3.3.3 Gaussian Low-Pass Filters

A Gaussian low-pass filter is defined as

$$H(u, v) = \exp\left(-\frac{D^2(u, v)}{2D_0^2}\right),$$

where D_0 is the standard deviation of the Gaussian filter.

3.4 Image Sharpening

3.4.1 Ideal High-Pass Filters

An ideal high-pass filter passes all frequencies outside a circle of radius D_0 from the origin and suppresses all frequencies inside of this circle, without attenuation:

$$H(u, v) = \begin{cases} 0 & \text{if } D(u, v) \leq D_0, \\ 1 & \text{if } D(u, v) > D_0. \end{cases}$$

High-pass equivalent filters can be generated for all of the above low-pass filters by taking the complement of the low-pass filter:

$$H_{\text{HP}}(u, v) = 1 - H_{\text{LP}}(u, v).$$

3.4.2 Butterworth High-Pass Filters

A Butterworth high-pass filter is defined as

$$H(u, v) = \frac{1}{1 + (D_0/D(u, v))^{2n}},$$

where n is the order of the filter.

3.4.3 Gaussian High-Pass Filters

A Gaussian high-pass filter is defined as

$$H(u, v) = 1 - \exp\left(-\frac{D^2(u, v)}{2D_0^2}\right),$$

where D_0 is the standard deviation of the Gaussian filter.

3.4.4 Laplacian Filters

We can consider an alternative formulation for the Laplacian using a filter the frequency domain. In the frequency domain, the Laplacian becomes

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2} = -4\pi^2 u^2 F(u, v) - 4\pi^2 v^2 F(u, v) = -4\pi^2 (u^2 + v^2) F(u, v).$$

Let us then define the filter $H(u, v)$ as

$$H(u, v) = -4\pi^2 (u^2 + v^2),$$

or, with respect to the center of the frequency rectangle,

$$H(u, v) = -4\pi^2 \left((u - P/2)^2 + (v - Q/2)^2 \right) = -4\pi D^2(u, v).$$

Then, the Laplacian of an image $f(x, y)$ is given by

$$\nabla^2 f(x, y) = \mathcal{F}^{-1} \{ H(u, v) F(u, v) \}.$$

Then for the enhanced image

$$g(x, y) = f(x, y) + c \nabla^2 f(x, y),$$

$c = -1$, so that in the frequency domain, this is equivalent to

$$G(u, v) = F(u, v) - H(u, v) F(u, v).$$

3.5 Selective Filtering

Bandreject and bandpass filters can be used to selectively filter specific frequency bands in an image. A bandreject filter is defined as

$$H(u, v) = \begin{cases} 0 & \text{if } D_0 - \frac{W}{2} \leq D(u, v) \leq D_0 + \frac{W}{2}, \\ 1 & \text{otherwise,} \end{cases}$$

where D_0 is the centre frequency and W is the width of the band. A bandpass filter is defined as the complement of the bandreject filter:

$$H_{BP}(u, v) = 1 - H_{BR}(u, v).$$

3.5.1 Butterworth Bandreject Filters

A Butterworth bandreject filter is defined as

$$H(u, v) = \frac{1}{1 + \left(\frac{D(u, v)W}{D_0^2 - D(u, v)^2} \right)^{2n}},$$

where n is the order of the filter.

3.5.2 Gaussian Bandreject Filters

A Gaussian bandreject filter is defined as

$$H(u, v) = 1 - \exp\left(-\left(\frac{D^2(u, v) - D_0^2}{D(u, v)W}\right)^2\right),$$

where D_0 is the centre frequency and W is the width of the band.

3.5.3 Notch Filters

Notch filters are used to remove frequencies in a predefined neighbourhood of the frequency rectangle (rather than being centred at the origin). Zero-phase-shift filters must be symmetric about the origin, so a notch filter transfer function with center frequencies (u_0, v_0) must have a corresponding notch filter transfer function with center frequencies $(-u_0, -v_0)$. The transfer function of a notch filter is constructed by multiplying the transfer functions of two highpass filter transfer functions whose centers have been translated to the centers of the notches. The general form for a notch filter with Q notches is

$$H(u, v) = \prod_{k=1}^Q H_k(u, v) H_{-k}(u, v),$$

where $H_k(u, v)$ and $H_{-k}(u, v)$ are the transfer functions of the highpass filters centred at the k th notch and its negative, respectively.

4 Colour

Colour is a perceptual property of light that is produced when white light is passed through a prism. The wavelengths of light that are produced comprise the **electromagnetic spectrum**, in which the human eye can discern wavelengths between 400 nm and 700 nm.

4.1 Colour Fundamentals

The human eye contains three types of cones that can be roughly classified as being sensitive to three principal sensing regions:

- **Short-wavelength cones (S-cones):** 2% of cones that are sensitive to blue light (435.8 nm). These are also the most sensitive cones, despite only making up 2% of the total.
- **Medium-wavelength cones (M-cones):** 33% of cones that are sensitive to green light (546.1 nm).
- **Long-wavelength cones (L-cones):** 64% of cones that are sensitive to red light (700 nm).

These **primary** colours can be combined to produce the **secondary** colours of light:

- **Cyan:** A combination of blue and green light.
- **Magenta:** A combination of blue and red light.

- **Yellow:** A combination of green and red light.

Mixing the primary colours of light produces white light. The primary colours of **pigment** are the secondary colours of light, and mixing these produces black pigment. The primary colours are also referred to as the **additive primaries**, while the secondary colours are referred to as the **subtractive primaries**.

4.1.1 Characterising Colour

Colour can be characterised by three attributes:

- **Brightness:** The intensity of the colour.
- **Hue:** The dominant wavelength of light.
- **Saturation:** The relative purity of the colour, or the amount of white light mixed with a hue. Hence, the degree of saturation is inversely proportional to the amount of white light mixed with a hue.

Hue and saturation can also be combined to form the **chromaticity** of a colour.

4.2 Colour Models

Colour models, or color spaces, are coordinate system for representing colours mathematically. Some common colour spaces include:

- **RGB:** The most common colour space, which represents colours as a combination of red, green, and blue intensities. RGB is used in displays and cameras.
- **CMY(K):** A subtractive colour model that represents colours as a combination of cyan, magenta, and yellow intensities, with an optional black intensity. CMYK is used in colour printing.
- **HSI:** A cylindrical colour space that represents the hue, saturation, and intensity of a colour. This is also known as HSV (value) or HSL (lightness), where the intensity may be scaled to a different range.

4.2.1 RGB Colour Model

The RGB colour model can be represented as a unit cube, with a corner at the origin, with the primary colours of red, green, and blue at the vertices:

- (1, 0, 0): Red.
- (0, 1, 0): Green.
- (0, 0, 1): Blue.

The colour black is at the origin, while the colour white is at the opposite corner. Grayscale colours lie along the diagonal of the cube which intersects these two points. Note that in this model, colour values are normalised to the range $[0, 1]$.

4.2.2 CMY(K) Colour Model

The CMY(K) colour model exists on the same cube as the RGB colour model, but with the primary colours of cyan, magenta, and yellow at the vertices:

- (0, 1, 1): Cyan.
- (1, 0, 1): Magenta.
- (1, 1, 0): Yellow.

Mathematically, CMYK colours are the complement of RGB colours:

$$\begin{bmatrix} C \\ M \\ Y \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} - \begin{bmatrix} R \\ G \\ B \end{bmatrix}.$$

4.2.3 HSI Colour Model

The HSI colour model is a cylindrical colour space that represents colours as a combination of hue, saturation, and intensity. The hue value is the angle around this cylinder, with red at 0° , green at 120° , and blue at 240° . The intensity value is the distance along the intensity axis in the RGB cube, while the saturation value is the distance from the intensity axis.

$$\theta = \arccos \left(\frac{2R - G - B}{2\sqrt{(R - G)^2 + (R - B)(G - B)}} \right), \quad (1)$$

$$H = \begin{cases} \frac{\theta}{360^\circ} & \text{if } B \leq G, \\ 1 - \frac{\theta}{360^\circ} & \text{if } B > G, \end{cases} \quad (2)$$

$$S = 1 - \frac{3}{R + G + B} \min(R, G, B), \quad (3)$$

$$I = \frac{R + G + B}{3}. \quad (4)$$

To convert from HSI to RGB, we can use the following equations. In the RG sector ($0^\circ \leq H \leq 120^\circ$):

$$\begin{aligned} B &= I(1 - S), \\ R &= I \left(1 + \frac{S \cos H}{\cos(60^\circ - H)} \right), \\ G &= 3I - (R + B). \end{aligned}$$

In the GB sector ($120^\circ \leq H \leq 240^\circ$):

$$\begin{aligned} R &= I(1 - S), \\ G &= I \left(1 + \frac{S \cos(H - 120^\circ)}{\cos(60^\circ - (H - 120^\circ))} \right), \\ B &= 3I - (R + G). \end{aligned}$$

In the BR sector ($240^\circ \leq H \leq 360^\circ$):

$$\begin{aligned} G &= I(1 - S), \\ B &= I \left(1 + \frac{S \cos(H - 240^\circ)}{\cos(60^\circ - (H - 240^\circ))} \right), \\ R &= 3I - (G + B). \end{aligned}$$

4.3 Colour Image Acquisition

Colour images are acquired using monochromatic charge-coupled devices (CCDs) that are sensitive to red, green, and blue light. These devices are arranged in a Bayer filter mosaic, which consists of a 2×2 grid of sensors with one red, one blue, and two green sensors.

$$\begin{bmatrix} G & R \\ B & G \end{bmatrix}$$

Here two green sensors are used to improve the resolution of the green channel, which is the most sensitive to the human eye. When demosaicking the image, the colour channels are extracted from the mosaic, and missing entries are interpolated using the values of neighbouring pixels in that channel.

$$\begin{bmatrix} G_{11} & R_{12} \\ B_{21} & G_{22} \end{bmatrix} \rightarrow \begin{bmatrix} - & R_{12} \\ - & - \end{bmatrix} + \begin{bmatrix} G_{11} & - \\ - & G_{22} \end{bmatrix} + \begin{bmatrix} - & - \\ B_{21} & - \end{bmatrix}$$

4.3.1 Pixel Depth

The pixel depth of an image is the number of bits used to represent the colour of each pixel. The pixel depth determines the total number of colours that can be represented in an image. For an image with n bits per channel, the total number of colours is given by 2^{3n} . For example, an image with 8 bits per channel has a pixel depth of 24 bits and can represent $2^{24} = 16\,777\,216$ unique colours.

4.4 Colour Image Processing

Colour images can be processed using the same techniques as grayscale images, but with the additional consideration of the colour channels. Transformations can either be applied to each channel independently or to individual channels, such as the hue, saturation, and intensity channels in the HSI colour space. We can define colour image transformations in the same way as for grayscale images:

$$g(x, y) = T(f(x, y)),$$

where $f(x, y)$ is the input image, $g(x, y)$ is the output image, and T is the transformation function.

4.4.1 Colour Intensity Transformations

Intensity transforms take the form:

$$g(x, y) = kf(x, y),$$

- RGB: $s_i = kr_i$ for $i = 1, 2, 3$.
- HSI: $s_1 = r_1$, $s_2 = r_2$, and $s_3 = kr_3$.
- CMY: $s_i = kr_i + (1 - k)$ for $i = 1, 2, 3$.

4.4.2 Colour Complements

Colour complements are hues directly opposite each other on the colour circle. Complements are analogous to negative images in grayscale images and can be useful for enhancing detail in dark regions.

4.4.3 Colour Slicing

Colour slicing is a technique used to highlight specific colours in an image to separate objects from surroundings.

4.4.4 Colour Tonal Transformations

Colour tonal transformations are used to adjust the tonal range of an image. These transformations can be used to boost contrast in an image using power-law transformations.

4.4.5 Colour Balancing (White Balancing)

An image is colour imbalanced when a known white point in the image does not contain equal RGB or CMY components. Colour balancing is used to correct these imbalances by adjusting the channels to make such points truly white.

4.4.6 Colour Histogram Equalisation

Histogram equalisation can also be applied to the intensity channel of HSI images to improve the contrast of the image without altering the hue or saturation.

4.4.7 Colour Image Smoothing and Sharpening

Colour images can be smoothed or sharpened using the same techniques as grayscale images, by applying a filter to each channel independently. For example, an image can be smoothed by considering its K average neighbours:

$$y_{ij} = \frac{1}{K} \sum_{(s,t) \in S} x_{st},$$

where S is the set of K neighbours of pixel (i, j) . Similarly, an image can be sharpened by applying a Laplacian filter to each channel independently. The Laplacian of a colour image \mathbf{x} is simply

$$\nabla^2 \mathbf{x} = \begin{bmatrix} \nabla^2 x_1 \\ \nabla^2 x_2 \\ \nabla^2 x_3 \end{bmatrix}$$

4.4.8 Colour Image Segmentation

Colour image segmentation is the process of partitioning an image into regions based on colour. For HSI images, segmentation is performed on the hue channel, and the saturation channel is used to further isolate regions in the hue image. For RGB images, we can use distance metrics to measure if nearby pixels are similar to a desired colour **a**.